

Review

14 January 2019 20:34

Pipeline Hazards:

- Structural
 - Hardware cannot support a combination of instr
 - **Bubbles**
- Data
 - Instr depends on result from a instr still in the pipeline
 - **Forwarding**
 - **Stall 'Slots' - Compiler scheduling**
- Control
 - Delay between fetching instr and executing branches/jumps
 - **Early branch determination**
 - **SMT.**
 - **Get a free cycle to decode branch**
 - **While executing other thread**

Evaluating Pipeline designs:

- Influenced by
 - Compiler scheduling
 - Compiler instruction selection
 - Application behavior
- Influences
 - Number of instructions
 - Number of registers needed
 - Number of stalls
 - Clock time

Temporal Locality:

- if an item is referenced, it will tend to be referenced again soon

Spatial Locality:

- If an item is referenced, items whose addresses are close by tend to be referenced soon

Cache:

- Tag, Index, Block Offset
- DM, N-Way set, fully associative
- LRU vs ran replacement
- Write Through vs Write Back
 - WT buffer to L2\$
 - WB - dirty bit

Caches

18 January 2019 15:56

$$AMAT = HitTime + MissRate * MissPenalty$$

Reducing Miss Rate:

Misses:

- **Compulsory**
- **Capacity**
- **Conflict**

For same Capacity:

- **Block size**
 - initial improvement due to spatial locality
 - too large a block size wastes cache capacity - lower temporal
- **Associativity**
 - will increase hit rate, but increase hit time
 - higher AMAT as cache size increases
 - Solution
 - Way speculation
 - Victim Cache
 - For DM cache
 - access same time as L1
 - **very small**

Skewed-associative caches:

- use different indices in each way
 - **hash** index bits
- costs
 - latency of hash
 - LRU implementation

Hardware Prefetching

- Predict next data to be fetched
- Stream (FIFO) in parallel with cache
- On miss
 - check prefetch stream
 - avoid pollution
 - needs extra mem BW

Compiler:

- **software prefetching**
 - load to register or cache
 - cost of issuing
 - superscalar has more issue BW
 - often, HW just as effective
- **storage layout** transformations
 - Merging arrays
 - Multidimension array - match array layout to traversal order
- **iteration space** transformations
 - Loop interchange (nested order)
 - Loop fusion
 - Blocking - blocks of data vs rows/columns

Reducing Miss Penalty:

Write-Through - all the way

Write Back - only cache

Write Miss

- Allocate new cache line
 - read to fill block first
- Non-Allocate
 - Send to level data currently stored at

Large blocks - early restart / critical word first

Non-blocking cache:

- continue to supply hits

Add a second-level Cache:

- $global_{miss_rate} = L1_{miss_rate} * L2_{miss_rate}$
- different multi-level inclusion strategies

Reduce Cache Hit time:

- avoid address translation for L1\$
- virtually-indexed, physically-tagged
- limits a direct mapped cache to page-size
 - increase via associativity

Dynamic Scheduling

18 January 2019 15:56

Advantages:

- Handles cases when dependencies unknown
- Simplifies the compiler
- HW can run code intended for another pipeline

Issue Stage:

- **In-order**
- Begins Execution
- Split ID stage of 5-stage pipeline
 - Issue Stage
 - Decode
 - Check for structural hazards
 - Read Operands
 - Wait until no data hazards
 - Then read operands

Execution

- **Out-of-order Execution**
- Out-of-order **Completion**

Data Dependence:

- **True Dependence**
 - Flow of data
 - RAW
- Name Dependence
 - Can solve by **renaming**
 - No flow of data
 - **Anti-Dependence**
 - WAR
 - **Output Dependence**
 - WAW

Tomasulo's Algorithm:

- Goal
 - Increase effective number of registers
 - Renaming in hardware
- Associate register with a **changeable tag**

1. Issue - In Order
 - a. Feed **Reservation Station**
 - i. Source Operands
 - 1) Either Register Value
 - 2) Or Tag
 - ii. Free
 - 1) Avoid Structural Hazards
2. Execution - Out of Order
 - a. on Functional Units
 - b. Might have to wait for value from CDB
3. Write Back
 - a. May complete out-of-order
 - b. **Broadcast** result on **Common Data Bus**
 - i. With Tag
 - ii. Goes to RS and/or Registers

Precise Interrupts: - ROB

- Tomasulo's out-of-order completion
- Fix this with commit/retire stage
- **Reorder buffer**
 - Instruction, tag, value
 - In order
 - Commits to **commit-side registers**
- On Branch
 - If prediction correct
 - Continue
 - Else
 - ROB entries trashed
 - Issue-side registers reset from commit-side
- Misprediction penalty
 - Shorter pipeline
- Store-load forwarding in ROB

RUU: Register Update unit

- Combined RS and ROB
 - Tag - slot in RUU
 - Operands can wait on Tags
- Sits Between Dispatch and Scheduler
 - FUs write to RUU via CDB
- Connects to Commit

Branch Prediction

18 January 2019 15:57

1-bit BP Buffer

- lower bits of instruction address used as index
- taken / not-taken
- problem: in loop - 2 mispredictions

n-bit BP Buffer

- add *hysteresis*
- **2ⁿ states** - half take, half not-taken
- increasing number of entries or n will only take us so far
- bad for branches which **aren't highly biased** to taken or not-taken

Local History:

- instruction address

Global History:

- taken - not-taken history for all previously-executed branches
- Compromise - use **m** entries
 - m-bit Branch History Register
 - shift register

Correlating BP Buffer:

- **(m,n) "gselect"**
- use m bits to select one of 2^m n-bit BHTs
- real programs have low correlation

Tournament Branch Predictor

- 2 predictors: one global, one local
- combine with a **selector**
 - driven by a third predictor

Return Address Predictors

- function can be called from many locations
- hard to predict which address to return to
- save return address in a stack like buffer
 - 8-16 entries
 - catch jump for entries
 - check if correct prediction on return
 - entries trashes by interrupts/other programs

Warm-up:

- Simple predictors re-learn fast
- Sophisticated predictors re-learn more slowly
- **Selective** predictor chose simple until sophisticated warms up

Branch Target Buffer:

- indexed by lower PC bits
- if tag matches
 - value stored = predicted target address
- BTB can predict if current instruction will be a taken branch
 - without decoding instruction
- When a branch is *committed*
 - update BTB
 - value = target address
 - tag = upper branch instruction address
 - can have extra n-bit predictor for each entry
- if done before commit
 - can help with speculated instructions
- done after
 - BTB won't hold incorrect info

Predicated Execution:

- instead of branch
 - conditionally executed instructions
- reduce number of branches
 - reduce number of *mis-predicted* branches

Dependencies

18 January 2019 15:57

There are four types:

- Data ("true") dependence: $S1 \delta S2$
 - $OUT(S1) \cap IN(S2)$
- Anti dependence: $S1 \bar{\delta} S2$
 - $IN(S1) \cap OUT(S2)$
- Output dependence: $S1 \delta^o S2$
 - $OUT(S1) \cap OUT(S2)$
- Control dependence: $S1 \delta^c S2$

21:10

("S1 must write something before S2 can read it")

("S1 must read something before S2 overwrites it")

("If S1 and S2 might both write to a location, S2 must write after S1")

Statement — S_X
-instance S_X^I

Dependence from earlier iteration: $S_X^I \delta_{<} S_X^J$

Loop-carried:

True dep: $\delta_{<}$

Anti-dep: $\bar{\delta}_{<}$

Respect order: δ_*

Same iteration of S: $\delta_{=}$

Parallel: EPIC/VLIW

18 January 2019 15:57

Side-channels / Vulnerabilities

09 March 2019 23:10

Parallel: GPU

18 January 2019 15:57

Vectors

09 March 2019

23:12

Multicore & Clusters

18 January 2019 15:57