

# Anonymization

27 October 2019 13:29

## Pseudonymization:

- secret function to map **direct identifiers**
- can approximate polynomials
- use cryptographic hash instead
  - can be broken using lookup table
    - with guessing of hash function
    - length of original direct id
  - assume hash function known:
  - add salt to original direct identifier
  - **Salt:** fixed string of arbitrary length (but long!) that is added to the identifier before hashing it -must be kept a secret

**Sensitive information:** trying to protect

**Identifier:** directly identifies a person

## **Quasi-identifier:**

- does not directly identify a person
- multiple taken together could uniquely identify a person

**Auxiliary information:** information known to an attacker

## **Uniqueness w.r.t A:**

- fraction of the dataset that is uniquely identified by the set of A of quasi-identifiers

## **k-anonymous:**

- every record in the table is indistinguishable from at least k-1 other records, with respect to every set of quasi identifiers

## **Equivalence class:**

- set of records that have the same values for all the quasi-identifiers

## Achieving k-anonymity:

- Non-perturbative methods
  - generalization - replace attribute values with more general ones
    - e.g. 43221 -> 4322\* -> 43\*\*\*
  - Suppression: delete a column or row
- Perturbative methods
  - add noise
  - data swapping
    - swap attributes between individuals

## **Homogeneity attacks:**

- individuals in the same equivalence class all have the same sensitive attribute value
- prevention:  $\ell$  - Diversity
  - an equivalence class is  $\ell$ -diverse if it contains at least  $\ell$  distinct values for the sensitive attributes
  - a table is  $\ell$ -diverse if every equivalence class is  $\ell$ -diverse

## **Semantic attacks:**

- sensitive attributes of individuals in an equivalence class are distinct but *semantically* similar

## **Skewness attacks:**

- the distribution of sensitive attributes in a class is skewed

## **t-closeness:**

- **distance** between the distribution of a sensitive attribute in the equivalence class and the distribution of this attribute in the whole table is not more than a threshold **t**
- table has **t-closeness** if all classes have t-closeness

# Big Data Anonymization

27 October 2019 13:30

- no sensitive attribute
- no quasi-identifier

## Unicity $\epsilon_p$ :

- average fraction of the users in the dataset that are uniquely identified by  $p$  random points

## Estimating unicity:

- Random set on  $N$  users
- for  $u$  in users
  - draw  $p$  points at random
  - is the trace unique?

## Reducing with generalization:

- coarsen the data by reducing resolution
- helps, but is not sufficient

## Matching Attacks:

- auxiliary info might not directly match data
  - noise, missing, multiple matches
- between two datasets
  - Anonymized dataset
  - Dataset with direct identifiers
- Rely on
  - measure of *distance* between two points
  - *linking algorithm*

## Matching attack: location data

- Assumption - number of actions a user performs in a region at a time interval:  $A(u, l, t) \sim Po(\lambda_{l,t})$
- Location - anonymized
- Actions - directly identifiable
- Step 1:
  - for each  $u_{DI}$  and  $v_{anon}$  compute a **score**
- Step 2:
  - Compute **max weight matching** between  $U$  and  $V$  users, using Hungarian Algorithm
- Step 3:
  - An edge for user  $U_{DI}$  is only considered a match if it is **significant** ( $score > \epsilon * \sigma_{U_{DI} edges}$ )

## Profiling Attacks:

- Identifying users in an anonymous dataset using an identified dataset collected at a different time
- Step 1:
  - Extract a profile of the user in the identified dataset
  - through a profiling distance/algorithm
    - Jensen-Shannon divergence
    - Dot product
    - Cosine similarity
    - $L_1$
- Step 2:
  - Compare the profiles of known users to users in the anonymous dataset
  - to identify them using a linking algorithm
    1. Compute histograms in both datasets
    2. Compute the distance between each pair of histograms
    3. Use Hungarian algorithm to find the max weight matching
      - 1) same condition for 'good' match

# Query-Based Systems

07 November 2019 12:53

## QBS:

- Don't share the database.
- Provide aggregates for statistical purposes
  - e.g. counting queries
- Just QBS
  - Susceptible to uniqueness attacks

## Query Size Reduction:

- Query set must be over a certain threshold
- If not, don't return value
- Susceptible to **intersection** attacks

## Bounded Noise Addition:

- Perturb output of every query
- If adding non biased noise
  - $\tilde{A} = A + [-N, +N]$
  - $E(\tilde{A}) = A$
- Attacker Knows the bound
  - Calc diff between two queries
  - Make use of unique Aux info
  - Figure out  $A$  and  $A$
- Groups
  - If Group share a secret attribute
  - Can find out the value of the secret

## Unbounded:

- Solves above two issues
- Centered at 0 - don't introduce bias
- Bayes' Theorem
  - $P[A = x | \tilde{A} = 0]$   
$$= \frac{P[\tilde{A} = 0 | A = x] * P[A = x]}{P[\tilde{A} = 0]}$$
- Averaging Attacks
  - CLT
  - Bayes' - also gives posterior distribution
  - **Required number** of queries:
    - $n \geq 4\sigma^2 z_\alpha^2$
    - $z_{0.05} = 1.96$ ,  $z_{0.01} = 2.58$
  - Defend
    - Add consistent noise
    - Cache query - return cached value
    - Seeded PRNG
  - Get around defense
    - **Semantic averaging attacks**
    - Logically Equivalent Query but expressed differently

## Diffix:

- Sticky noise
  - For each condition
  - Add static
    - Random value
    - Seeded
      - $hash(C, salt)$
  - Add dynamic
    - Seed hash
      - Static seed
      - Unique ID of all conditions in the Query
- Bucket Suppression (QSR)
  - Dynamic, noisy threshold
  - If  $Q(D) \leq 1 \rightarrow suppress$
  - Else  $T \sim N(4, 0.5)$ 
    - Seeded hash
    - Salt and UID of C in Q
    - Suppress if  $Q(D) < T$
- Split Averaging Attacks
  - Pair of Semantically same query
  - Across Range of an attribute
  - Average these to get desired
- Diffix blocks by having static on each condition

# Differential Privacy

12 November 2019 13:28

$$1 - \epsilon \leq \frac{\Pr(\text{output} = y \mid x \in D)}{\Pr(\text{output} = y \mid x \notin D)} \leq 1 + \epsilon$$

## Neighboring Dataset:

- D1 neighbors D2 if they differ by only one row

Formal DP:

- $\Pr(M(D) = y) \leq e^\epsilon \Pr(M(D') = y)$ 
  - $e^{\pm\epsilon} \approx 1 \pm \epsilon$

Hard to protect **against group attacks** with DP

- Scale with group size ( **$k\epsilon$** )

Achieving DP:

- $f_X(x|\mu, b) = \frac{1}{2b} \exp\left(-\frac{|x - \mu|}{b}\right)$ 
  - $\sigma^2 = 2b^2$
  - Like a sharper Gaussian
- $Lap(b) = f_X(x|0, b)$
- For one query
  - **Add  $Lap(1/\epsilon)$**

Protect against Averaging attacks - **Composability:**

- Releasing output of any two queries protected by  $\epsilon - DP$
- Same as releasing one query protected by  $\epsilon - DP$

## Privacy Budget

- Add  **$Lap(1/\epsilon_i)$**  to  $Q_i$
- With the constraint:  $\epsilon = \sum_i^n \epsilon_i$

Optimizing DP: Histograms (counts):

- Add  $Lap(1/\epsilon)$  to every bucket, as opposed to multiple queries
  - Since a user only contributes to one bucket
  - Don't need to split budget

Optimizing DP Function Sensitivity:

- Global sensitivity of a function  $f$ 
  - $\Delta f = \max |f(D) - f(D')|$  (for any  $D$  and  $D'$ )
- Add noise  $Lap\left(\frac{\Delta f}{\epsilon}\right)$

Local DP:

- User gives **randomizes response**
- $\Pr(M(r_1) = r) \leq e^\epsilon \Pr(M(r_2) = r)$

# Computing on Encrypted Data

01 December 2019 20:49

## Secure Multi-Party Computation:

- Private inputs
- Jointly compute over a function
- One or more decrypt output

## Multiparty Addition Protocol

- Split shares
- Send to parties not named in share
- Broadcast **partial sums**

## Yao's Millionaire problem:

- $Z[i] = \text{Priv}_A(\text{Pub}_A(r_{\text{bsecret}} - b) + i) \bmod p_{\text{rand}}$
- $z[i] += 1 \text{ for } i > a$
- Z values all differ by at least 2 from each other
  - So that post increment they differ by at least 1

## **Honest but curious**

- Follows protocol but opportunistic (can collude)

## **Malicious**

- Deviates from protocol (can collude)

## Asynchronous Communication

- Results sent asynchronously

## **Fair**

- Secure against not forwarding last message

## Oblivious Transfer Protocol (1 from 2):

- A → B Two public keys
- B chooses one, and send symmetric K
- A decrypts twice, sends  $K(m1), K_{\text{bad}}(m2)$ 
  - Other way round if B chose 2nd pubA
- Bob decrypts, with K, same choice

## Yao Garbled Gates (B doesn't know gate function)

- A create two random keys for each wire
- A computes truth table, sends **shuffled** one to B
- A sends choice of A keys to B, (A inputs)
- Using an **OT**, **B selects a KBX for each wire**
- A tells B final K mappings

- $k[w, 0]p[w]$  encrypts a 0
- $k[w, 1]$  **not**  $p[w]$  encrypts a 1

- Permute bits index the table
- Perform gate on encoded values
- Selects new permute

- B just index's and is told what keys to use to decrypt

## One Time Memory:

- Store secret K0 and K1
- Select with input
- Third bit, x, says if used

## One Time Program:

- Convert to Yao Garbled Circuit
- Use OTM to store k[]'s

## Extending Yao to two circuits:

- Want to compute  $f_B(a, b)$  as well as  $f_A(a, b)$
- B computes
- $f(a, b, k) = k \oplus f_A(a, b) \oplus f_B(a, b)$
- B sends first part to A
- A decrypts by XOR'ing with secret bit k

# Computing on Untrusted Servers

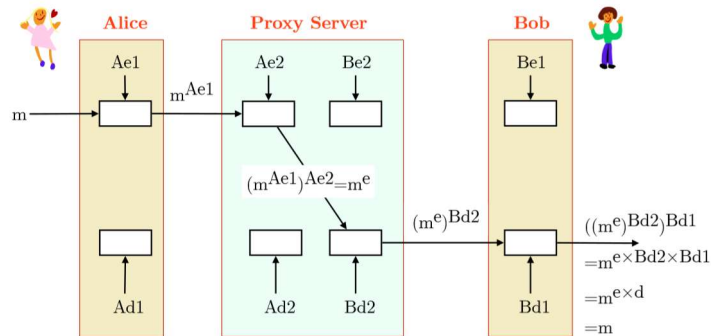
03 December 2019 15:43

## CryptDB:

- encrypt query
- Query made on encrypted data
- results decrypted
- **Onions** of Encryptions
  - multiple layers of encryption
  - CryptDB has different onions for different operators
  - an **attribute might have multiple** onioned values depending on application
- based on secret key shared by clients
  - (decrypt results)
- hard to remove a user

## Proxy-based encryption:

- Key server generates a master RSA key pair
  - $(e, n), (d, n)$
- For each user
  - a pair for user  $(e1, n), (d1, n)$
  - a pair sent to DB server  $(e2, n), (d2, n)$
  - $e1 * e2 = e \text{ mod } (p-1)(q-1)$
  - $d1 * d2 = d \text{ mod } (p-1)(q-1)$



## Symmetric Bilinear Pairings:

- pair two elements from one group to an element of a second group
- $G_1, G_2$  are two cyclic groups of prime order  $q$
- $g$  is a generator for  $G_1$
- bilinear pairing  $G_1 * G_1 \rightarrow G_2$ 
  - for all  $u, v \in G_1, a, b \in \mathbb{Z}_q$
  - $e(u^a, v^b) = e(u, v)^{ab}$
  - $e(g, g) \neq 1$
  - $e(u, v)$  computable

# Anonymous Communication I

03 December 2019 15:44

## Chaum Mixes (Mix-Nets):

- Mixer T - trusted
  - Forwards to A
  - $K_t(a, K_a(m))$
  - reorders outbound forward
    - adds delay
- Listener knows number of inward and outbound, but can't match them
- Improve by doing  $A \rightarrow t_1 \rightarrow t_2 \rightarrow B$ 
  - $K_{t_1}(t_2, K_{t_2}(b, K_b(m)))$
- Anonymous Reply
  - $K_t(b, K_b(m, K_t(a, K_x), K_y))$
  - $K_x, K_y$  one time pair, or symmetric

## Onion Routing:

- relay nodes
- e.g.  $K_{r_1}(r_2, K_{r_2}(\dots(B, m)))$

## Tor:

- clients creates virtual circuits with **relay nodes**
- Inter relay encrypted using TLS
- 512 byte cell size

## Controls Cells

- CID(2): Circuit ID for the link
- CMD(1):
  - Create/Destroy circuit
- PAYLOAD(509):
  - additional control data

## Relay Cells:

- CID(2)
- CMD(1) = Relay
- Next encrypted with AES session key for this link (made during control)
  - MISC(10)
  - CMD(1)
    - extend/extended
      - sends new AES keys for next relay back to A
    - begin/connected (TCP connection)
    - Data
  - Payload(498)

## Circuit Construction:

- P is tor proxy for A
- P -> R1 Create
- R1 -> P Created (AES K1)
- P -> R1 Relay, Extend R2
- R1 -> R2 Create
- R2 -> R1 Created (AES K2)
- R1 -> P Relay, Extended (K2)

## TCP

- P -> R1 Relay (R2, Begin B)
- R1 -> R2 Relay (Begin B)
- R2 -> R1 TCP handshake B
- R2 -> R1 Relay Connected
- R1 -> P Relay Connected

## Location-Hidden Services:

- B creates Tor circuit to **Introduction Relays**
  - Inform of **Long-term Public** key
- Publish Name, Public Key, and IR's to **lookup service**
- A finds this then
  - Asks via **IR** for B to connect to A via
    - a *rendezvous relay V* (separate from IRs)

# Anonymous Communication II

03 December 2019 23:57

## Attacks against Tor:

- **Global traffic analysis**
  - fingerprint B/W over time
- **Active interference attacks**
  - congest victim to relay
  - monitor which relay to public server flow is affected
  - Overcome - **Collective Control Plane (CCP)**
    - DC-nets for secret inputs to public outputs
    - managed by **CCP Policy Oracle**
      - control when/ how much to send
- **Denial-of-service**
  - on lookup service
- **Intersection attacks**
  - Intersects with users in non Tor that match quasi-identifiers
  - That also used Tor at a certain time
  - **Buddies**
    - Policy Oracle reports metrics on simulated intersection attacks
    - Possinymity: possibilistic deniability
      - intersection of users for a Nym for a time period
    - Indinymity: probabilistic indistinguishability

## DC-Nets:

- Dining Cryptographers
  - **XOR** on left, right, + your secret
  - Send to middle - another **XOR**
  - Doesn't scale well
  - If 1 user leaves, have to restart
  - 1 Malicious user can jam communication
- 
- **De-anonymizing exploits**
    - attack the browser
  - **Accountability provisions**
    - Before Unmasking as last resort
      - threat of censure
      - give an opportunity to retract
      - expulsion from group



# Privacy Policies

04 December 2019 10:56

## S4P Language:

- assumes user trust service providers to enforce user policies
- supports policy evolution

## User:

- may assertion -> gives permission
- will query -> asks for promise

## Service:

- may query -> asks for permission
- will assertion -> gives promise

For satisfactions, **compare** mays and wills (**queries to assertions**)

Both a **user privacy preference** and **service privacy policy** consist of a set of assertions and a query.

## Assertion

$E \text{ says } f_0 \text{ if } f_1, \dots, f_n \text{ where } c$

Delegation of authority

## Fact

$f ::= a \mid e \text{ may } b \mid e \text{ will } b \mid e \text{ can say } f$

## Query

$q ::= e \text{ says } f? \mid c? \mid \neg q \mid q_1 \wedge q_2 \mid q_1 \vee q_2 \mid \text{exists } x: q$

<b>E</b>	<b>constant</b>	Typically <b>principals</b> e.g. Alice, Bob, Service Provider
<b>x</b>	<b>variable</b>	
<b>e</b>	<b>expression</b>	Constant or variable
<b>c</b>	<b>constraint</b>	Constraint on variables occurring in assertion
<b>a</b>	<b>atom</b>	<b>Predicate</b> written in <i>infix</i> notation e.g. 'Alice is a nicePerson'
<b>b</b>	<b>behaviour atom</b>	<b>Service behaviours</b> e.g. 'delete email within 1 yr'